



TAM Service Engineer Engagement Check

Puppet Enterprise Architecture Assessment Report

for

Acme Corp

Prepared By: Adam Shambrook

Last Revision: v1.0.0

Last Revision Date: February 1, 2022

Document Summary

This Technical Account Management Service Engineer (TSE) quarterly assessment check report has been prepared after obtaining Puppet Server and Puppet Agent diagnostic and log information and focusing on Architecture specific items. The Puppet TSE and Solution Architect teams have reviewed your specific environment details against Puppet Patterns and Tactics associated with Puppet Enterprise implementation architecture.

This report will be presented during a customer advisory session to review the noted findings, suggested Puppet environment changes / updates and other Puppet recommendations. Optionally, your Puppet account team will follow-up with a SOW for a Puppet Professional Services Engineer or Puppet Solutions Architect to implement the recommendations if your team is unable to perform the suggested changes.

Engagement Check Contacts

	Puppet TAM	Puppet TSE	Puppet Architect
Name:	[Insert TAM Name]	[Insert TSE Name]	[Insert Architect Name]
Email:	[Insert TAM Email]	[Insert TSE Email]	[Insert Architect Email]

Engagement Check Summary

The Puppet Enterprise Architecture assessment check will review your current Puppet Enterprise architecture to validate against supported Puppet patterns, assess current server sizing against near and long term planned growth and assess your current configuration against the recommended configuration.

Recommendations

The following recommendations are suggested based on the Architecture assessment check:

1. Current server sizing and load:
Puppet Enterprise Architecture: Large
Node Count: 16,000
Primary Server Size: 25 CPU, 32 GB RAM
CPU resource allocation: Total/Used/Free 25/16/9
2. Guidance for future scaling based on business needs:
 - a. Your Puppet Primary server has 9 CPUs more than our highest recommended allocation. This is over-dimensioned, so could be scaled down, and the additional CPUs allocated elsewhere.
 - b. At 16,500 nodes, and at your current rate of growth, we estimate that you will reach 20,000 nodes within the next 12 months. For 20,000+ nodes, we recommend the Extra Large Puppet Enterprise Architecture. You should consider moving to that architecture to

ensure sufficient performance. This installation is managed by Puppet's [PEADM module](#), which will ensure your architecture is configured correctly automatically.

- c. Your Hiera configuration looks correct, however if you experience a difference in performance of PE across regions, consider whether an alternate one of Puppet's [multi-region reference architectures](#) might improve performance.
3. PE tuning suggestions for efficiency:
 - a. No other applications should be hosted on the Primary than the Puppet-related services installed by default. To ensure Puppet has access to its anticipated resources, and improve performance, remove additional applications that are currently hosted on your Primary server.
4. Other Noted Items:
 - a. Your Puppet Replica node appears to be smaller than your primary server, at 8 CPU and 16GB RAM. This is not sufficient to restore your primary server in the event of a failover and needs to be increased to match the size of your Primary Server.
 - b. With an estate of your size, consider CD4PE to manage testing and the deployment of your Puppet Enterprise code.